

Introduction to Machine Learning

Chapter 15: Bagging and Random Forests

Bernd Bischl, Christoph Molnar

Department of Statistics – LMU Munich

Winter term 2017/18



ENSEMBLE METHODS

- Ensemble methods combine the predictions of several base learners and combine them into an aggregated estimator.
- Homogeneous ensembles (multiple models of same base learner)
 - Bagging: Fit models on bootstrapped versions of training data
 - Boosting: runs sequentially: each model on reweighted data version / the previous residuals so it improves the errors of the previous round
- Heterogeneous ensembles (different base learners)
 - Fit different base learners on the same data or different “views” of the same data. Then learn how to aggregate their predictions, often with a 2nd-layer model.

ENSEMBLE METHODS

General homogenous approach (often it works like this but not always)

- A “base learner” is selected and fitted multiple times to either resampled or reweighted versions of the original data.
- The base learner is applied to either resampled or reweighted versions of the original dataset. This results in M prediction functions $b^{[1]}(x), \dots, b^{[M]}(x)$.
- These M function are aggregated, usually in a linear fashion. This results in the following final prediction function:

$$f(x) = \sum_{m=1}^M \beta^{[m]} b^{[m]}(x)$$

with coefficients $\beta^{[1]}, \dots, \beta^{[M]}$.

BAGGING

- Bagging is based on **B**ootstrap **A**ggregation.
- Proposed by Breiman (1996).
- Train on multiple bootstrap samples of data \mathcal{D} , then combine:
 - ❶ Create M bootstrap samples of size n .
 - ❷ Fit the base learner on each of the M bootstrap samples.
 - ❸ Aggregate the predictions of the M estimators via averaging or majority voting.
- M affects Monte-Carlo approximation error; main hyperparameter.
- Interpretability of the model becomes harder.

BAGGING

Bagging algorithm

- 1: **Input:** Dataset \mathcal{D} , base learner, number of bootstraps M
- 2: **for** $m = 1 \rightarrow M$ **do**
- 3: Draw a bootstrap sample $\mathcal{D}^{[m]}$ from \mathcal{D} .
- 4: Train base learner on $\mathcal{D}^{[m]}$, obtain model $b^{[m]}(x)$
- 5: **end for**
- 6: Aggregate the predictions of the M estimators (via averaging or majority voting), to determine the bagging estimator:

$$f(x) = \frac{1}{M} \sum_{m=1}^M b^{[m]}(x)$$

BAGGING

- Bagging reduces the variance of the estimator, but increases the bias in return.
- Bagging works best for unstable/high variance learners (learners where small perturbations in training set lead to larger changes in the prediction)
 - Classification and regression trees
 - Neural networks
 - Piece-wise variable selection in the regression case, etc.
- For stable estimation methods bagging might degrade performance
 - k-nearest neighbor
 - discriminant analysis
 - naive bayes
 - linear regression

WHY DOES BAGGING WORK?

- Suppose we have a numerical dependent variable.
- The training datasets are given by \mathcal{D} and base learner estimator by $f(x)$.
- The datasets are sampled independently from distribution \mathbb{P}_{xy} (data generating process).
- The *theoretical* aggregated estimator is given by

$$f_A(x) = \mathbb{E}_{\mathcal{D}}[f(x)].$$

WHY DOES BAGGING WORK?

- Let x, y be a random sample from \mathbb{P}_{xy} but independent of \mathcal{D} . The average error of the normal $f(x)$ is then $e = E_{\mathcal{D}} E_{xy}[(y - f(x))^2]$ and of the aggregated estimator $e_A = E_{xy}[(y - f_A(x))^2]$.
- It follows:

$$e = E_{\mathcal{D}} E_{xy}[(y - f(x))^2] = E_{xy}[y^2] - 2E_{xy}[yf_A] + E_{\mathcal{D}} E_{xy}[f^2(x)]$$

- And we apply Jensen's inequality to e :

$$\begin{aligned} e = E_{xy} E_{\mathcal{D}}[(y - f(x))^2] &\geq E_{xy}(E_{\mathcal{D}}[y - f(x)])^2 = E_{xy}[(y - f_A(x))]^2 = e_A \\ &= E_{xy}[y^2] - 2E_{xy}[yf_A] + E_{xy}[f_A^2(x)] \end{aligned}$$

WHY DOES BAGGING WORK?

- The difference between e and e_A is $E_{\mathcal{D}}E_{xy}[f^2(x)] \geq \mathbb{E}_{xy}[f_A^2(x)]$
- The more unstable $f(x)$, the more error reduction we obtain.
- But the bagging estimator only approximates the theoretical f_A (bootstrap), we therefore suffer from approximation error (bias) by using the empirical distribution function instead of the true data generating process and only perform M bootstrap iterations instead of an infinite number.
- Bagging does not necessarily lead to an improved classifier.
 - Example: binary outcome, $y = 1$ for all values of x
 - Consider random classifier f with $P(f(x) = 1) = 0.4$ (independent of x)
 - Prediction error for f is 0.6
 - Prediction error for the bagging estimator is 1

RANDOM FORESTS

- Modification of bagging for trees.
- Proposed by Breiman (2001).
- Construction of bootstrapped **decorrelated** trees
- The variance of the bagging prediction depends on the correlation between the trees ρ

$$\rho\sigma^2 + \frac{1-\rho}{M}\sigma^2,$$

where σ^2 describes the variance of a tree.

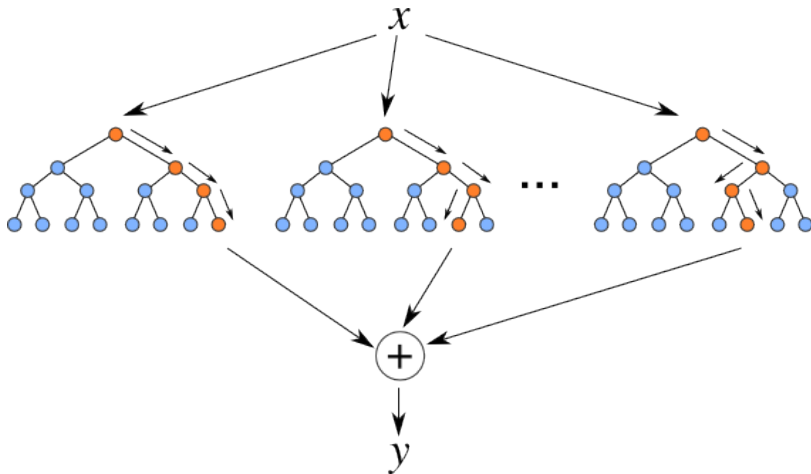
- ⇒ Reduce correlation by randomization in each split. Instead of all p features, draw $m_{try} \leq p$ random split candidates.
- ⇒ Trees are expanded as much as possible, without aggressive early stopping or pruning, to increase variance.

RANDOM FORESTS

Random Forest algorithm

- 1: **Input:** A dataset \mathcal{D} of n observations, number M of trees in the forest, number $mtry$ of variables to draw for each split
 - 2: **for** $m = 1 \rightarrow M$ **do**
 - 3: Draw a bootstrap sample $\mathcal{D}^{[m]}$ from \mathcal{D}
 - 4: Grow tree $b^{[m]}(x)$ using $\mathcal{D}^{[m]}$
 - 5: For each split only consider $mtry$ randomly selected features
 - 6: Grow tree without early stopping or pruning
 - 7: **end for**
 - 8: Aggregate the predictions of the M estimators (via averaging or majority voting), to predict on new data.
-

RANDOM FORESTS



RANDOM FORESTS

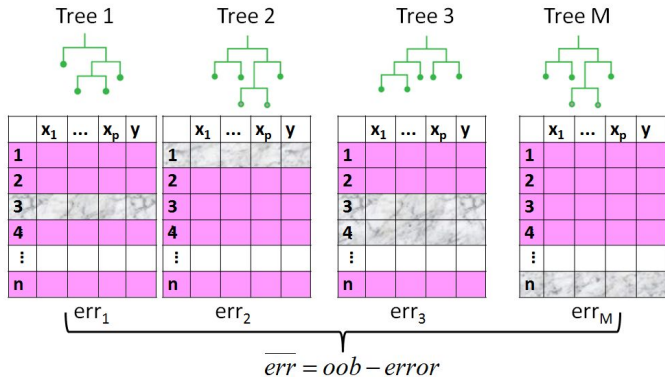
- The following values are recommended for *mtry*:
 - Classification: $\lfloor \sqrt{p} \rfloor$
 - Regression: $\lfloor p/3 \rfloor$
- Out-Of-Bag error: On average ca. 1/3 of points are not drawn.

$$\mathbb{P}(\text{Obs. not drawn}) = \left(1 - \frac{1}{n}\right)^n \xrightarrow{n \rightarrow \infty} \frac{1}{e} \approx 0.37.$$


To compute the OOB error, each observation x is predicted only with those trees that did not use x in their fit.


- The OOB error is similar to cross-validation estimation. It can also be used for a quicker model selection.

RANDOM FORESTS



$$err_m = \sum_{x^{(i)} \in OOB_m} L(y^{(i)}, \hat{y}_m^{(i)})$$

 in-bag observations, used to build the trees (Remember: The same observation can enter the in-bag sample more than once.)

 out-of-bag observations (OOB_m), used to evaluate prediction performance (err_m)